KNN - Tutorial

Hisham Ihshaish
Bristol, UK
June, 2021

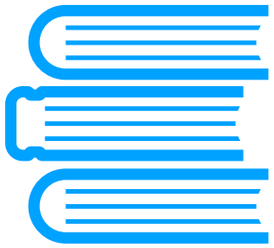

Do we evaluate our ‘model’ manually?

Nah!

lecture 2

Evaluate,

in training and in testing

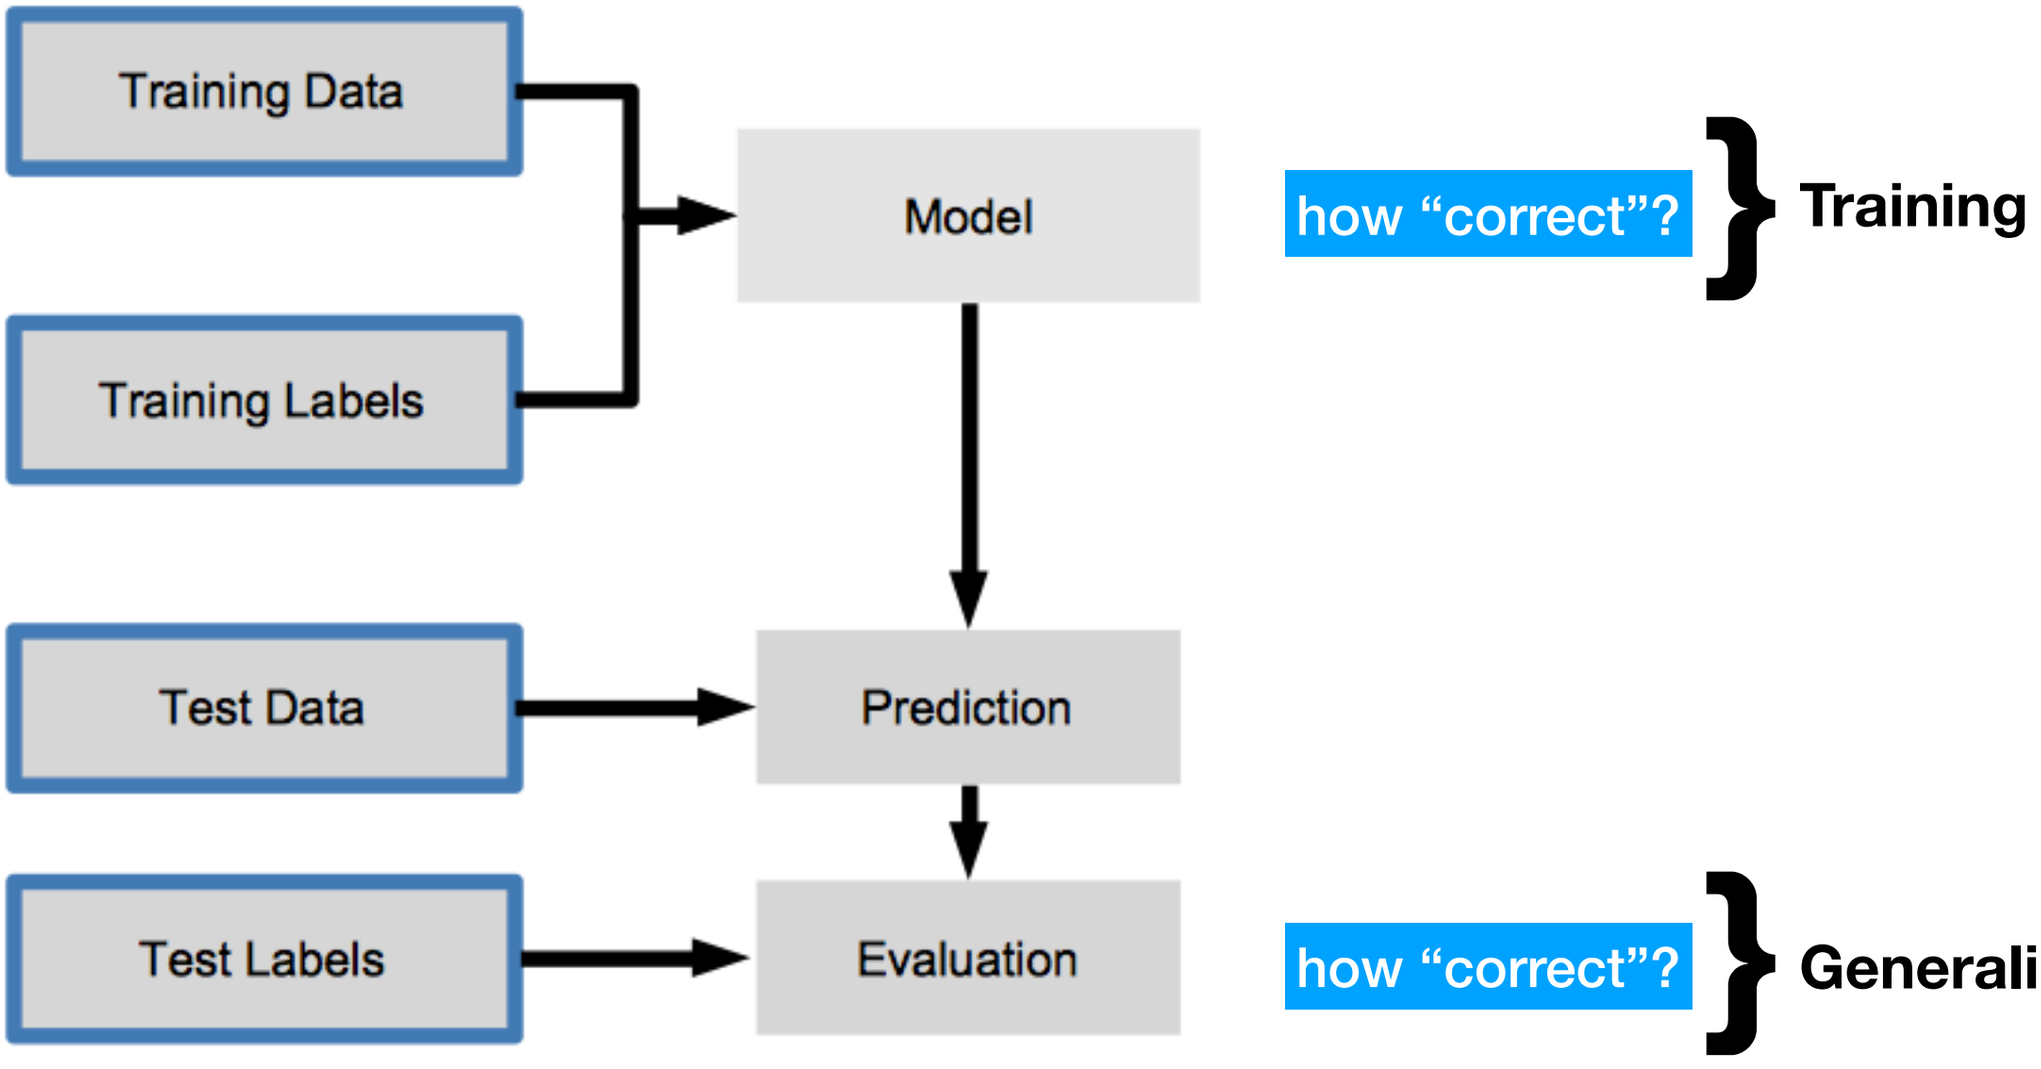

Data… splitting

- for **training** and for **testing** (generalisation)

$$X = \begin{pmatrix} 1.1 & 2.2 & 3.4 & 5.6 & 1.0 \\ 6.7 & 0.5 & 0.4 & 2.6 & 1.6 \\ 2.4 & 9.3 & 7.3 & 6.4 & 2.8 \\ 1.5 & 0.0 & 4.3 & 8.3 & 3.4 \\ 0.5 & 3.5 & 8.1 & 3.6 & 4.6 \\ 5.1 & 9.7 & 3.5 & 7.9 & 5.1 \\ 3.7 & 7.8 & 2.6 & 3.2 & 6.3 \end{pmatrix} \quad y = \begin{pmatrix} 1.6 \\ 2.7 \\ 4.4 \\ 0.5 \\ 0.2 \\ 5.6 \\ 6.7 \end{pmatrix}$$

Data… splitting

- for **training** and for **testing** (generalisation)

Fitting the model

1.1	2.2	3.4	5.6	1.0
6.7	0.5	0.4	2.6	1.6
2.4	9.3	7.3	6.4	2.8
1.5	0.0	4.3	8.3	3.4
0.5	3.5	8.1	3.6	4.6
5.1	9.7	3.5	7.9	5.1
3.7	7.8	2.6	3.2	6.3

1.6
2.7
4.4
0.5
0.2
5.6
6.7

Training Data — **~ 70% - 80%**

Testing Data — **~ 20% - 30%**

Generalisation error, or say, model accuracy!

Run example - explained

```
import numpy as np
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
import matplotlib.pyplot as plt
%matplotlib inline
```

```
df = pd.read_csv('/Users/test/Downloads/breast-cancer-wisconsin.data.txt')
```

read data

```
df.replace('?', -99999, inplace=True)
df.drop(['id'], 1, inplace=True)
```

preprocessing

```
X = np.array(df.drop(['class'], 1))
y = np.array(df['class'])
```

identify X and y

```
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)


clf = KNeighborsClassifier(n_neighbors=1)
clf.fit (X_train, y_train)

print (clf.score(X_train, y_train))
print (clf.score(X_test, y_test))

data = np.array([4,3,3,2,1,2,1,1,2])

prediction= clf.predict(data.reshape(1,-1))

print(prediction)
```

$$X = \begin{pmatrix} 1.1 & 2.2 & 3.4 & 5.6 & 1.0 \\ 6.7 & 0.5 & 0.4 & 2.6 & 1.6 \\ 2.4 & 9.3 & 7.3 & 6.4 & 2.8 \\ 1.5 & 0.0 & 4.3 & 8.3 & 3.4 \\ 0.5 & 3.5 & 8.1 & 3.6 & 4.6 \\ 5.1 & 9.7 & 3.5 & 7.9 & 5.1 \\ 3.7 & 7.8 & 2.6 & 3.2 & 6.3 \end{pmatrix} \quad y = \begin{pmatrix} 1.6 \\ 2.7 \\ 4.4 \\ 0.5 \\ 0.2 \\ 5.6 \\ 6.7 \end{pmatrix}$$

Run example - explained

```
import numpy as np
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
import matplotlib.pyplot as plt
%matplotlib inline
```

```
df = pd.read_csv('/Users/test/Downloads/breast-cancer-wisconsin.data.txt')
```

read data

```
df.replace('?', -99999, inplace=True)
df.drop(['id'], 1, inplace=True)
```

preprocessing

```
X = np.array(df.drop(['class'], 1))
y = np.array(df['class'])
```

$$X = \begin{pmatrix} 1.1 & 2.2 & 3.4 & 5.6 & 1.0 \\ 6.7 & 0.5 & 0.4 & 2.6 & 1.6 \\ 2.4 & 9.3 & 7.3 & 6.4 & 2.8 \\ 1.5 & 0.0 & 4.3 & 8.3 & 3.4 \\ 0.5 & 3.5 & 8.1 & 3.6 & 4.6 \\ 5.1 & 9.7 & 3.5 & 7.9 & 5.1 \\ 3.7 & 7.8 & 2.6 & 3.2 & 6.3 \end{pmatrix} \quad y = \begin{pmatrix} 1.6 \\ 2.7 \\ 4.4 \\ 0.5 \\ 0.2 \\ 5.6 \\ 6.7 \end{pmatrix}$$

```
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

Split into training data and testing data

```
clf = KNeighborsClassifier(n_neighbors=1)
clf.fit (X_train, y_train)

print (clf.score(X_train, y_train))
print (clf.score(X_test, y_test))

data = np.array([4,3,3,2,1,2,1,1,2])

prediction= clf.predict(data.reshape(1,-1))

print(prediction)
```

Data… splitting

- for **training** and for **testing** (generalisation)

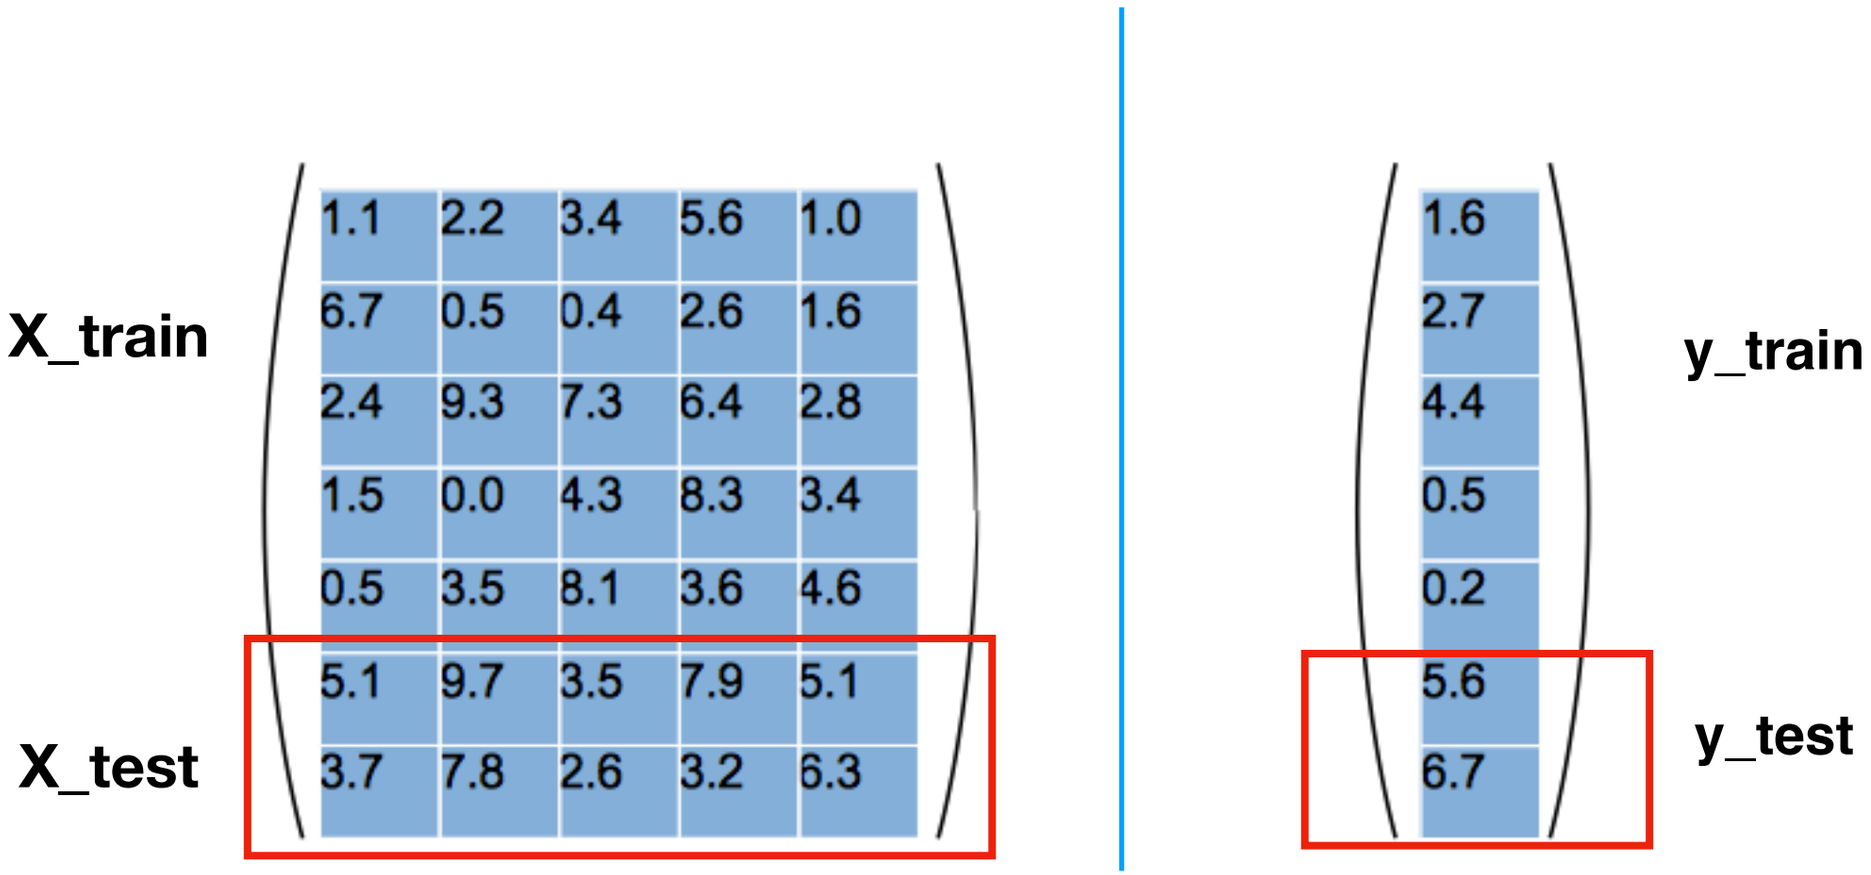

sklearn >> **X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)**

here goes what we did last week

```
import numpy as np
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
import matplotlib.pyplot as plt
%matplotlib inline
```

```
df = pd.read_csv('/Users/test/Downloads/breast-cancer-wisconsin.data.txt')
```

→ ***read data***

```
df.replace('?', -99999, inplace=True)
df.drop(['id'], 1, inplace=True)
```

→ ***preprocessing***

```
X = np.array(df.drop(['class'], 1))
y = np.array(df['class'])
```

→ X = | 1.1 2.2 3.4 5.6 1.0 |
| 6.7 0.5 0.4 2.6 1.6 |
| 2.4 9.3 7.3 6.4 2.8 |
| 1.5 0.0 4.3 8.3 3.4 |
| 0.5 3.5 8.1 3.6 4.6 |
| 5.1 9.7 3.5 7.9 5.1 |
| 3.7 7.8 2.6 3.2 6.3 |

y = | 1.6 |
| 2.7 |
| 4.4 |
| 0.5 |
| 0.2 |
| 5.6 |
| 6.7 |

```
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

Split into training data and testing data

```
clf = KNeighborsClassifier(n_neighbors=1)
clf.fit (X_train, y_train)
```

train, only on <u>training data</u>

```
print (clf.score(X_train, y_train))
print (clf.score(X_test, y_test))
```

→ ***training error***

→ ***generalisation error***

```
data = np.array([4,3,3,2,1,2,1,1,2])

prediction= clf.predict(data.reshape(1,-1))

print(prediction)
```

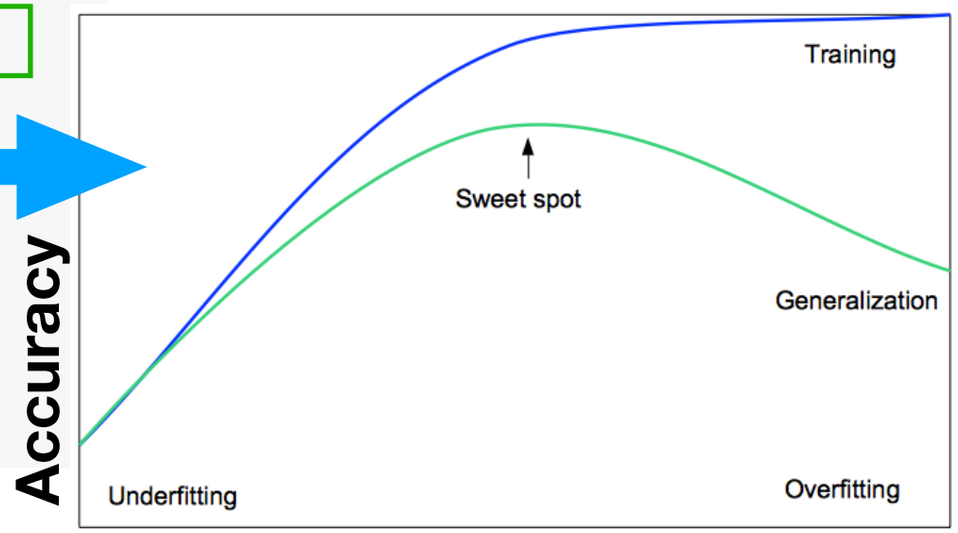

note on file upload

Preparation:

- Call libraries
- Read file:

a) if on **local jupyter notebook** simply use `df = pd.read_csv('/Users/test/Downloads/breast-cancer-wisconsin.data.txt')`

b) if in colabd, import io, then upload file as shown below.

- Encapculate file in a Dataframe — (preferred) check methods of df to show first rows, shape, and statistics of your data.

[31]
```
import numpy as np
import pandas as pd
import io

from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
import matplotlib.pyplot as plt


#df = pd.read_csv('/Users/test/Downloads/breast-cancer-wisconsin.data.txt')

from google.colab import files
uploaded = files.upload()


df = pd.read_csv(io.BytesIO(uploaded['breast-cancer-wisconsin.data.txt']))
```