

A: Simple Linear Regression

Diabetes Dataset - can be loaded from sklearn (mean centred and scaled - read here: <http://scikit-learn.org/stable/datasets/index.html>). Original data can be found in <https://www4.stat.ncsu.edu/~boos/var.select/diabetes.html>

For Regression, one feature out of 10 provided is selected, namely the *bmi* (body mass index). The response of interest (*y*) is a quantitative measure of disease progression one year after baseline, which, hypothetically is linearly dependent on *bmi*.

Tasks

- 1) Plot *bmi* and *y*
- 2) Execute provided code and discuss:
 - a) the coefficients for the linear model (intercept and slope).
 - b) the Coefficient of Determination (R^2).
 - c) is the linear model provided a good predictive model, why?

```
1 # print(__doc__)
2 # Code source: Jaques Grobler
3 # License: BSD 3 clause
4 # changes made for teaching purpose by Hisham Ihshaish
5
6 import pandas as pd
7 import matplotlib.pyplot as plt
8 import numpy as np
9 from sklearn import datasets
10 from sklearn.metrics import mean_squared_error, r2_score
11
12 from sklearn.model_selection import train_test_split
13 from sklearn.linear_model import LinearRegression
14
15 # Load the diabetes dataset
16 diabetes = datasets.load_diabetes()
17 #if you want to display dataset keys uncomment below.
18 #print (diabetes.keys())
19
20 data = pd.DataFrame(diabetes.data, columns=[diabetes.feature_names])
21 target = pd.DataFrame(diabetes.target)
22 #print(data.describe())
23 X = diabetes.data[:, np.newaxis, 2]
24 y = np.array(target)
25
26 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)
27 lr = LinearRegression().fit(X_train, y_train)
28
29 print("lr.coef_: {}".format(lr.coef_))
30 print("lr.intercept_: {}".format(lr.intercept_))
31
32 # Make predictions using the testing set
33 y_pred = lr.predict(X_test)
34
35 print('Coefficient of Determination R Squared: %.2f' % r2_score(y_test, y_pred))
36
37 # Plot outputs
38 plt.scatter(X_test, y_test, color='black')
39 plt.plot(X_test, y_pred, color='blue', linewidth=3)
40 plt.xticks(())
41 plt.yticks(())
42 plt.show()
```

B1: Multiple Linear Regression

Boston House Prices Dataset - can be loaded from sklearn (http://scikit-learn.org/stable/modules/generated/sklearn.datasets.load_boston.html). Original data can be explored in <https://www.cs.toronto.edu/~delve/data/boston/bostonDetail.html>

I extended the number of features, originally 14, from the features themselves - using a method 'PolynomialFeatures' provided sklearn (may be out of scope here, but the main idea is that we want the high-dimensional features' space, ie, number of independent variables is large = 104 produced). The response of interest (y) is an estimate for House Price, which is, supposedly, linearly dependent on the features provided.

Tasks

- 1) Execute the provided code and discuss
 - a) what is the test score metric?
 - b) why scores are different on training and on testing datasets?
 - c) what's the reason behind the produced difference?

```
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import numpy as np
4 from sklearn import datasets
5 from sklearn.preprocessing import MinMaxScaler, PolynomialFeatures
6 from sklearn.model_selection import train_test_split
7 from sklearn.linear_model import LinearRegression
8
9 #boston dataset extended: added features using PolynomialFeatures
10 def load_extended_boston():
11     boston = datasets.load_boston()
12     X = boston.data
13
14     X = MinMaxScaler().fit_transform(boston.data)
15     X = PolynomialFeatures(degree=2, include_bias=False).fit_transform(X)
16     return X, boston.target
17
18 X, y = load_extended_boston()
19
20 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
21 lr = LinearRegression().fit(X_train, y_train)
22 print("Training set score: {:.2f}".format(lr.score(X_train, y_train)))
23 print("Test set score: {:.2f}".format(lr.score(X_test, y_test)))
24
```

B2: Multiple Linear Regression

Use Ridge Regression and tune alpha hyper-parameter.

Tasks

- 1) Produce a scatterplot matrix for the set of features/attributes - use **pandas.plotting.scatter_matrix** class (check examples and documentation).
- 2) Can you produce a VIF table for your independent variables — you can use: `from statsmodels.stats.outliers_influence import variance_inflation_factor`
- 3) Explore potential multicollinearity between the attributes — how would you design your linear regression model if evident multicollinearity exists.
- 4) Execute the provided code and discuss
 - a) Are scores provided by both regression models different? Explain why?
 - b) Tune alpha hyper-parameter and explore prediction scores - explain results.
 - c) Produce a residuals plot for each case — you either implement a calculation yourself or use yellowbrick's class — check [here](#).

```
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import numpy as np
4 from sklearn import datasets
5 from sklearn.preprocessing import MinMaxScaler, PolynomialFeatures
6 from sklearn.model_selection import train_test_split
7 from sklearn.linear_model import LinearRegression
8 from sklearn.linear_model import Ridge
9
10
11 #boston dataset extended: added features using PolynomialFeatures
12 def load_extended_boston():
13     boston = datasets.load_boston()
14     X = boston.data
15
16     X = MinMaxScaler().fit_transform(boston.data)
17     X = PolynomialFeatures(degree=2, include_bias=False).fit_transform(X)
18     return X, boston.target
19
20 X, y = load_extended_boston()
21
22 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
23 lr = LinearRegression().fit(X_train, y_train)
24
25 ridge = Ridge(alpha=1.0).fit(X_train, y_train)
26
27 print("Training set score: {:.2f}".format(lr.score(X_train, y_train)))
28 print("Test set score: {:.2f}".format(lr.score(X_test, y_test)))
29 print("Ridge Training set score: {:.2f}".format(ridge.score(X_train, y_train)))
30 print("Ridge Test set score: {:.2f}".format(ridge.score(X_test, y_test)))
31
```

C: Lasso Regression

Tasks

```
>> lasso = Lasso().fit(X_train, y_train)
```

Apply Lasso Regression and tune its alpha hyper-parameter: compare with the scores you obtained earlier.

B3: Produce Figure in lecture notes - slide no 33

```
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import numpy as np
4 from sklearn import datasets
5 from sklearn.preprocessing import MinMaxScaler, PolynomialFeatures
6 from sklearn.model_selection import train_test_split
7 from sklearn.linear_model import LinearRegression
8 from sklearn.linear_model import Ridge
9
10
11 #boston dataset extended: added features using PolynomialFeatures
12 def load_extended_boston():
13     boston = datasets.load_boston()
14     X = boston.data
15
16     X = MinMaxScaler().fit_transform(boston.data)
17     X = PolynomialFeatures(degree=2, include_bias=False).fit_transform(X)
18     return X, boston.target
19
20 X, y = load_extended_boston()
21
22 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
23 lr = LinearRegression().fit(X_train, y_train)
24
25 ridge = Ridge(alpha=1.0).fit(X_train, y_train)
26 ridge10 = Ridge(alpha=10).fit(X_train, y_train)
27 ridge01 = Ridge(alpha=0.1).fit(X_train, y_train)
28 plt.plot(ridge.coef_, 's', label="Ridge alpha=1")
29 plt.plot(ridge10.coef_, '^', label="Ridge alpha=10")
30 plt.plot(ridge01.coef_, 'v', label="Ridge alpha=0.1")
31
32 plt.plot(lr.coef_, 'o', label="LinearRegression")
33 plt.xlabel("Coefficient index")
34 plt.ylabel("Coefficient magnitude")
35 xlims = plt.xlim()
36 plt.hlines(0, xlims[0], xlims[1])
37 plt.xlim(xlims)
38 plt.ylim(-25, 25)
39 plt.legend()
40 plt.show()
```